


Raytracing: Quality

COSC 4328/5327
Scott A. King

Ray Genealogy

- Primary rays spawn off 3 rays. Two of those can spawn off 3 more, etc.
- When do you stop?
 - When ray leaves the scene.
 - When the contribution is small enough.




After each bounce the contribution is attenuated by the k 's in the illumination model. Check this value and set a maximum recursion level.

Ray Genealogy

- 1 Ray/pixel at $1k \times 1k$ image = 1M rays.
- Say on avg. 6 secondary rays = 7M rays
- 100k objects with 10 ops for intersection = 7,000,000M ops
- 4GHz processor 5 cycles per op = 800MFLOPS
- How long to render? $70000/8 = 8750s \sim 146m \sim 2.5hr$
- Will this image look really good for our 2.5hr?

Describe what you see!




(Hint)




©Zachary Wartell - 1/26/2005

Describe what you see!



©Zachary Wartell - 1/26/2005


Describe what you see!



Aliasing = a high frequency signal masquerading as a low frequency signal due to a poor sampling frequency.

©Zachary Wartell - 1/26/2005

Recall: frequency, cycle/wavelength



$$y = A \sin(x / \Delta x_{\text{cycle}})$$

A is amplitude

Δx_{cycle} is wavelength

(ignoring y-offset and phase shift)

$$f = 1 / \Delta x_{\text{cycle}} \quad \text{frequency}$$

What is a "poor sampling frequency"?

Nyquist Sampling Frequency:

$$f_s = 2 f_{\text{max}}$$


$$\Delta x_s = \Delta x_{\text{cycle}} / 2$$

f_s : Nyquist Sampling Frequency

f_{max} : maximum frequency component in signal

Δx_s : Nyquist Sampling Interval


Δx_{cycle} : Cycle interval



"To avoid losing information we need to set sampling frequency to at least twice the highest frequency component of signal."

Example of highest frequency component

- What is highest frequency component of this signal, $S(x)$?




TADA!

- What if I tell you the equation is: $\sin[x (2 \pi / 10)] + \sin[x (2 \pi / 1)] * 0.25$ comp. of $S(x)$

$$\sin[x (2 \pi / 10)] + \sin[x (2 \pi / 1)] * 0.25$$

$$f = 1/10, a = 1$$


$$f = 1, a = 0.25$$




$$\sin[x (2 \pi / 10)] + \sin[x (2 \pi / 1)] * 0.25$$

Super-Sampling

- Split single pixel into sub-pixels.
- Pixel's final color is a mixture of sub-pixels' colors.
Simple method: Sample at the middle of each sub-pixel.
Then, pixel's color is the average of the sub-pixels' color.

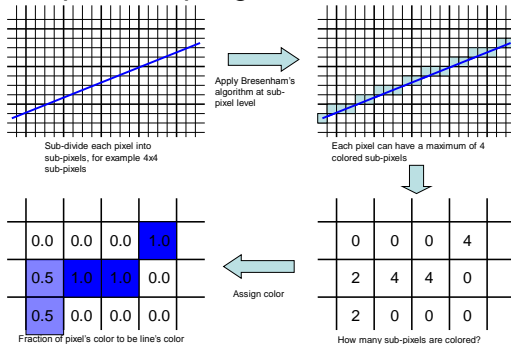


Pixels




Sub-divide into sub-pixels

Super-Sampling a Zero-Width Line



Super-Sampling a Line with Non-Zero Width



Sub-pixel Weighting Masks


- Instead of considering each sub-pixel to be of equal importance, assign a weight to each sub-pixel.
- Usually consider the center sub-pixel to be most important

	1	2	2	1
	2	3	3	2
	2	3	3	2
	1	2	2	1

Example weight for each sub-pixel
Total weight = 32

Final color of pixel = Sum of each (sub-pixel color x sub-pixel weight) / total weight

Common Filter Functions



$$\text{Gaussian function} = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$


where σ is the standard deviation and μ is the mean

Super-Sampling

- A simple method is to shoot more than one ray per pixel and average the values.


Super-Sampling

- If we shoot one ray/pixel we get one sample/pixel
- What if we shoot more rays/pixel?
- How many rays to double the resolution?
- How about triple the resolution?
- How many rays/pixel to remove aliasing?




Super-Sampling

- How many total rays for an NxN image?




Adaptive Supersampling




Adaptive Supersampling

- Tries to only fire rays that are needed
- Is the assumption that if the rays are the same color, that color is the right pixel color valid?
 - Small objects cause problems
 - Slender objects cause problems
 - If moving very noticeable artifacts



Distributed/Stochastic Raytracing


- We can't remove all aliasing.
- What can we do to lessen how noticeable the artifacts are?
- How does the news stop you from recognizing a person?
- Change the aliasing into blurring
- Do that by randomizing the sampling frequency.



Monte Carlo Integration

- Let rectangle g enclose a function f .
- Choose set of random points within g .
- How does the # of points within g , n_g , compare with the # of points within f , n_f ?
- That is, what is

$$\frac{n_f}{n_g}(b-a)c$$




- As n_g increases what happens?

Monte Carlo Integration

- Let's use random numbers more efficiently.
- Instead of picking (x, y) pairs, just pick random x values and plug them into the function.
- How do we use the y values for integration?
- The average y times the interval gives an estimate of area.

$$\int_a^b f(x) dx \approx \frac{b-a}{n} \sum_{i=1}^n f(x_i)$$



Monte Carlo Integration

- In general

$$I \approx \frac{V}{n} \sum_{i=1}^n f(x_i)$$

Where V is the measure of the set to integrate over. So if we had

$$\int_{a_1}^{b_1} \int_{a_2}^{b_2} \dots \int_{a_d}^{b_d} f(x_1, x_2, \dots, x_d) dx_1 dx_2 \dots dx_d$$

then

$$V = \prod_{i=1}^d (b_i - a_i)$$

For example:


$$\int_1^3 \int_0^5 \int_0^1 \int_0^2 f(x, y, z, t) dx dy dz dy \approx \frac{24}{n} \sum_{i=1}^n f(x, y, z, t)$$

Random Numbers

- What makes a sequence random?
- A sequence is **uniformly distributed** if all subsets contain a random sequence. For example, say we have random points in a plane. The number of points that fall in a rectangle with area k should only depend on k and not the location of the rectangle.
- In a computer, random numbers are generated *deterministically* and are not truly random. They are often called *pseudorandom*.
- Not all random number generators are equal. And in fact most are down right bad (at least those that come with compilers.)

Pseudo vs. Quasi

- Pseudo a bit too random for Monte Carlo, use Quasi instead
 - For super-sampling, generate evenly spaced samples then jitter within sub-pixel.



Motion Blur



Source: wikipedia

Motion Blur

- Sample the image temporally; temporal anti-aliasing.
- Can perform with other rendering techniques.
- Modeling the shutter of a camera.
- In ray tracing, can combine with spatial anti-aliasing by giving each ray a jittered time.





Cook, Porter, Carpenter 1984




Depth of Field

- Model the lens of a camera.
- All points in the scene project as a circle on the image plane, called the *circle of confusion*.
 - Objects at the focal distance are sharp, others are blurred.
- Not only jitter the ray direction, jitter where on the lens it hits




Soft Shadows

- Hard shadows result from using point light sources.
- Instead, model a light source with some geometry.
- If only part of the light is visible, determine that amount and decrease light accordingly.
- Stochastic shadow rays work well.



Ravi Ramamoorthi



Links

- <http://www.povray.org> (POV)
- <http://www.irtc.org> (IRTC)
- <http://radsite.lbl.gov/radiance/HOME.html> (radiance)
- <http://www-graphics.stanford.edu/~cek/rayshade/rayshade.html> (rayshade)
- <http://www.acm.org/tog/resources/RTNews/html/> (ray tracing news)
- <http://www.cs.cmu.edu/~efros/java/tracer/tracer.html>